

Computer Science 2XC3 - Final Project

Due Date: April 9, 2025 WINTER

Group 9

Wu, Peter: `wu729`

Ganesh, Vikram: `vikramgv`

Bai, Franklin: `baif4`

Contents

Part 1, Team Charter	2
Part 2, Single Source Shortest Path Algorithms	2
Part 3, All-pair Shortest Path Algorithm	2
Part 4, A* Algorithm	2
Part 5, Compare Shortest Path Algorithms	2
Part 6, Organize code per UML Diagram	2

Part 1, Team Charter

Communication

For our group, we decided that communication over **Discord** would be the most convenient for all members. Discord also allows screen sharing, making it a valuable tool for collaboration. Each member is expected to respond within **one hour** to maintain efficiency and work within our given timeframe.

For planned video calls, members are expected to join within **15 minutes** of the scheduled time. Since these meetings take up a significant portion of our schedules, punctuality is appreciated. If a member is unable to join, they must notify the group beforehand so the rest can plan accordingly. Any health or family emergencies will, of course, be waived.

Repeated failure to adhere to the above communication agreements will result in an email to a TA or the professor, outlining the member's inability to meet agreed timeframes. This may lead to a deduction in their grade.

Collaboration Tools

Our team will use **GitHub** to maintain version control. Each member may use the IDE of their choice, provided it is properly synced with Git. A shared repository will be created to store all code, diagrams, and other essential resources. This allows us to collaborate simultaneously, monitor each other's progress, and assist remotely.

The final report will be written using **LateX** to take advantage of its formatting capabilities compared to other programs. Sections may be drafted in other software and later copied or converted into LateX.

Dispute Resolution

Team disputes will be resolved via a discussion on Discord or through a scheduled video call. During this discussion, the team will review the reasons behind any disagreements. If necessary, a second chance may be offered.

If disputes cannot be resolved within the team and repeated failures occur, the matter will be escalated to a TA or the professor via email.

Deadlines and Task Allocation

Our first internal deadline is **March 27th**, during which we will have a video call to discuss our progress on individual components of the project. We will have another internal deadline around **April 7th**, 2 days before the deadline, where we will begin finalizing our project.

Initial task distribution is as follows:

- **Franklin:** Part 1, 6
- **Vik:** Part 2, 3
- **Peter:** Part 4, 5
- **Team:** Part 7

These assignments are not final. Members are encouraged to assist one another, as we recognize the difficulty level may vary between different parts of the project.

Part 2, Single Source Shortest Path Algorithms

In section 2, the modified versions of the dijkstra and bellman ford algorithms are implemented, where each algorithm was allowed to perform at most k relaxations per node. Each implementation was then tested and compared against the original versions of the algorithms for accuracy, performance and memory usage. The results of the tests are shown in the tables below.

Accuracy

To measure the accuracy of the modified algorithms, we designed various experiments to determine how the size of the graph, density of the graph, and the number of relaxations per node (k value) affected the accuracy of the algorithms. The results of these experiments are shown in the graphs below.

The first experiment was designed to measure the accuracy of the modified algorithms with different k values. The results of this experiment are shown in Figure 1.




Figure 1: Experiment 1 - Accuracy of modified algorithms with different k values.

This experiment was performed across 50 trials, where each trial consisted of a graph with 100 nodes, 200 edges and a maximum weight of 10. The results of this experiment show that the accuracy of the modified Bellman-Ford algorithm increases as the k value increases. This is because the modified Bellman-Ford algorithm is allowed to perform at most k relaxations per node, which means that they may not be able to find the shortest path in the graph if the k value is too small.

However, the value of k does not seem to have any effect on the accuracy of the modified Dijkstra's algorithm, and the modified Dijkstra's algorithm is able to find the shortest path in the graph regardless of the value of k. This is because the modified Dijkstra's algorithm is able to find the shortest path in the graph by using a priority queue to keep track of the nodes that have been visited and the nodes that have not been visited. Therefore, the value of k does not affect the accuracy of the modified Dijkstra's algorithm.

The second and third experiments were designed to measure the accuracy of the

modified algorithms with different graph sizes and densities. The results of these experiments are shown in Figure 2 and Figure 3.




Figure 2: Experiment 2 - Accuracy of modified algorithms with different graph sizes.




Figure 3: Experiment 3 - Accuracy of modified algorithms with different graph densities.

Experiment 2 was performed across 50 trials, where each trial consisted of a graph with $k = 10$, a maximum edge weight of 10 and a density D of 0.5. The experiment was conducted with $|V| = 10, \dots, 100$. The number of vertices were calculated based on the number of edges in the graph using the following formula:

$$\frac{|E|}{2} = |V|(|V| - 1)0.5$$

where $|E|$ is the number of edges, $|V|$ is the number of vertices and D is the density of the graph.

Experiment 3 was performed across 50 trials, where each trial consisted of a graph with $k = 10$, a maximum edge weight of 10 and 20 vertices. The experiment was conducted with $D = 0.1, 0.2, \dots, 1.0$. The density of the graph was calculated based on the number of edges in the graph using the following formula:

$$\frac{|E|}{2} = 20(20 - 1)D$$

where $|E|$ is the number of edges, $|V|$ is the number of vertices and D is the density of the graph.

The results of these experiments show that the accuracy of the modified Bellman-Ford and dijkstra's algorithms are not affected by the size of the graph or the density of the graph. We can see this in Figure 2 and Figure 3, where the accuracy of the modified Bellman-Ford and dijkstra's algorithms are 100% for all graph sizes and densities.

The results of these experiments show that the accuracy of the modified Bellman-Ford and dijkstra's algorithms are not affected by the size of the graph or the density of the graph, but Bellman-ford alone is affected by the number of relaxations allowed per node.

Time Complexity

To measure the time complexity/performance of the modified algorithms, we designed various experiments to determine how the size of the graphs, density of the graphs, and the number of relaxations per node (k value) affected the performance

of the algorithms. The results of these experiments are shown below.

Experiment 4 was designed to measure the performance of the modified algorithms with different k values. This experiment was conducted with the same parameters as experiment 1, where each trial consisted of a graph with 100 nodes, 200 edges, a maximum edge weight of 10 and k values of 1, ..., 10. The results of this experiment are shown in Figure 4.




Figure 4: Experiment 4 - Performance of modified algorithms with different k values.

The results of this experiment varied for each algorithm. The modified Bellman-Ford algorithm was significantly faster than the original Bellman-Ford algorithm. However, the speed of the modified Bellman-Ford algorithm did increase as the value of k increased, with more of a linear increase. This shows that restricting the number of relaxations per node allows the Bellman-Ford algorithm to perform significantly faster. Furthermore, the modified dijkstra's algorithm performed almost identically to the original dijkstra's algorithm, showing that the modified dijkstra's algorithm is not affected by the value of k . Both versions of dijkstra's algorithm performed faster than both versions of the Bellman-Ford algorithm, which is expected as dijkstra's algorithm is more efficient and has a better time complexity than Bellman-Ford

($O(E + V \log V)$ vs $O(VE)$).

Experiment 5 was designed to measure the performance of the modified algorithms with different graph sizes. This experiment was conducted with the same parameters as experiment 2, where each trial consisted of a graph with $k = 10$, a maximum edge weight of 10 and a density D of 0.5. The number of vertices were calculated based on the number of edges in the graph using the same formula as experiment 2. The results of this experiment are shown in Figure 5.




Figure 5: Experiment 5 - Performance of modified algorithms with different graph sizes.

Here, the results are similar to experiment 4, with the exception of the classical Bellman-Ford algorithm's performance. Both versions of Bellman-Ford increased in time complexity as the number of vertices increased, with the classic Bellman-Ford algorithm decreasing in performance at a much faster rate than the modified Bellman-Ford algorithm. This is because the modified Bellman-Ford algorithm is able to perform at most k relaxations per node, which means that it is able to perform faster than the classical Bellman-Ford algorithm. Both versions of Dijkstra's algorithm saw a decrease in performance as the number of vertices increased, but both performed faster than each version of Bellman-Ford. The modified Dijkstra's

algorithm performed very similarly to the classical dijkstra's algorithm, which is expected as the modified dijkstra's algorithm is not affected by the value of k . This is to be expected, since dijkstra's algorithm is more efficient and has a better time complexity than Bellman-Ford ($O(E + V \log V)$ vs $O(VE)$).

Experiment 6 was designed to measure the performance of the modified algorithms with different graph densities. It was conducted with the same parameters as experiment 3, where each trial consisted of a graph with $k = 10$, a maximum edge weight of 10 and 20 vertices. The density of the graph was calculated based on the number of edges in the graph using the same formula as experiment 3. The results of this experiment are shown in Figure 6.




Figure 6: Experiment 6 - Performance of modified algorithms with different graph densities.

Here, we can see that the density of the graph affects the performance of all the algorithms in a similar way. Both versions of Bellman-Ford and dijkstra's algorithms saw a decrease in performance as the density of the graph increased, with the classical Bellman-Ford algorithm decreasing in performance at a much faster rate than the modified Bellman-Ford algorithm. Similar to experiment 5, this is because the restriction of k relaxations per node allows the modified Bellman-Ford algorithm to perform faster than the classical Bellman-Ford algorithm. The modified dijkstra's algorithm also performed similarly to the classical dijkstra's algorithm, however, the modified dijkstra's algorithm performed slightly slower than the classical dijkstra's algorithm. however, the difference in performance was very small. Both of these algorithms performed faster than both versions of Bellman-Ford, which is expected as dijkstra's algorithm is more efficient and has a better time complexity than Bellman-Ford.

Therefore, we can see that each parameter has different effects on the performance of each algorithm.

Memory Usage

Similar to the previous experiments, to calculate the memory usage of the modified algorithms, we designed various experiments to determine how the size of the graphs, density of the graphs, and the number of relaxations per node (k value) affected the memory usage of the algorithms.

Experiment 7 has designed to measure the memory usage of the modified algorithms with different k values. The parameters of the experiment were the same as experiment 1 and experiment 4, where each trial consisted of a graph with 100 nodes, 200 edges, a maximum edge weight of 10 and k values of 1, ..., 10. The results of this experiment are shown in Figure 7.




Figure 7: Experiment 7 - Memory usage of modified algorithms with different k values.

Experiment 7 shows that the memory usage of every algorithm does not seem to be affected by the value of k , as the memory usage of each algorithm is relatively the same for all values of k . This is expected as the memory usage of each algorithm is determined by the number of vertices and edges in the graph, rather than the value of k . Both the modified version of the Bellman-Ford algorithm and the modified version of dijkstra's algorithm used very similar amounts of memory as the classical versions of the algorithms. Both of the modified algorithms also use more memory than the classical versions of the algorithms, which is expected as the modified algorithms are more complex and have more variables to store and modify.

Therefore, we can see that the memory usage of each algorithm is only affected by the number of vertices, and the density and size of the graph do not affect the memory usage of each algorithm.

Experiment 8 was designed to measure the memory usage of the modified algorithms with different graph sizes. The parameters of the experiment were the same

as experiment 2 and experiment 5, where each trial consisted of a graph with $k = 10$, a maximum edge weight of 10 and a density D of 0.5. The number of vertices were calculated based on the number of edges in the graph using the same formula as experiment 2. The results of this experiment are shown in Figure 8.




Figure 8: Experiment 8 - Memory usage of modified algorithms with different graph sizes.

Experiment 8 shows that the memory usage of each algorithm increases as the number of vertices in the graph increases. The classical versions of both algorithms perform faster than the modified versions of the algorithms and use less memory than the modified versions of the algorithms. However, there is a clear trend among all algorithms, where the memory usage of each algorithm increases as the number of vertices in the graph increases. This is expected as the number of vertices in the graph increases, as we know that the space complexity of both algorithms is $O(V)$, where V is the number of vertices in the graph.

Experiment 9 was designed to measure the memory usage of the modified algorithms with different graph densities. The parameters of the experiment were the

same as experiment 3 and experiment 6, where each trial consisted of a graph with $k = 10$, a maximum edge weight of 10 and 20 vertices. The density of the graph was calculated based on the number of edges in the graph using the same formula as experiment 3. The results of this experiment are shown in Figure 9.

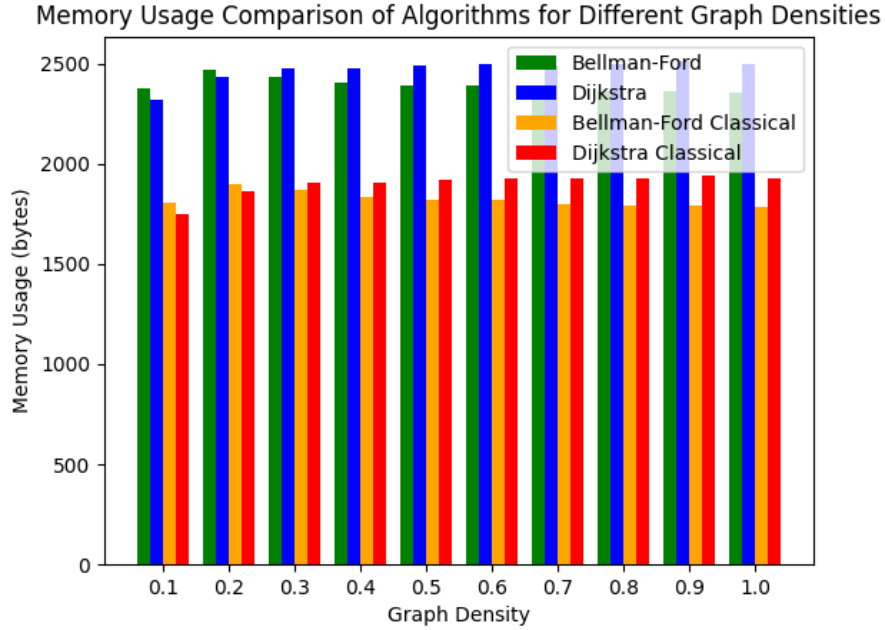


Figure 9: Experiment 9 - Memory usage of modified algorithms with different graph sizes.

Experiment 9 shows that the memory usage of each algorithm is not affected by the density of the graph, as the memory usage of each algorithm is relatively the same for all values of density. This is expected as the memory usage of each algorithm is determined by the number of vertices and edges in the graph, rather than the density of the graph. Hence, the results of this experiment are very similar to the results of experiment 7, where both the modified version of the Bellman-Ford algorithm and the modified version of Dijkstra's algorithm used very similar amounts of memory as the classical versions of the algorithms.

Part 3, All-pair Shortest Path Algorithm

In section 3, an algorithm to find the shortest path in a graph from all nodes to a single node were implemented. This was implemented in two different ways, one using dijkstra's algorithm and the other using bellman ford.

For dense graphs, the single-source implementation of dijkstra's algorithm has a time complexity of $O(V^2)$, where V is the number of vertices in the graph, while it's time complexity for sparse graphs is $O(E + V \log V)$, where E is the number of edges in the graph. The implementation of the shortest path algorithm (for all nodes) using dijkstras has a time complexity of $O(V^3)$ for dense graphs and $O(E^2 + EV^2 \log V)$. This is because dijkstra's algorithm is run V times, once for each node in the graph.

The same logic applies to the implementation of the shortest path algorithm using bellman ford. The time complexity for dense graphs is $O(V^3)$ and for sparse graphs is $O(E^2 + EV^2)$. The implementation of the shortest path algorithm using bellman ford has a time complexity of $O(V^4)$ for dense graphs and $O(E^2 + EV^2)$ for sparse graphs. This is because bellman ford is run V times, once for each node in the graph.

Part 4, A* Algorithm

Part 4.1

Issues A_star is trying to address

A_star algorithm is trying to address Dijkstra's algorithm performance problem in a very dense and complex graph, where it has full exploration of the entire graph that is computation heavy. Dijkstra's algorithm also doesn't have the idea of a source going with the destination, which means it would explore a vast number of unnecessary nodes, in the cases when we are not interested in other paths. On the other hand, A_star employs the heuristic function to guide the search towards more promising nodes. This provides an estimate of the cost from the current node to the goal node, allowing A_star to prioritize nodes closer to the goal and cut the unnecessary exploration.

Experiment for Dijkstra's vs A_star

Step 1: Graph Generation

Generate various random graphs with different levels of complexity. Create a set of graphs with the same number of nodes but varying densities. Also, generate another set of graphs with the same number of edges but different sizes (i.e., number of nodes).

Step 2: Keep Track of the Running Time

For each run, ensure that both algorithms are tested on the same graph, with the same randomly selected source node (*src*) and a randomly selected destination node (*dst*) for A_star. After each execution of the algorithm, record the running time.

Step 3: Graph Plotting

We should plot two graphs: both with the y-axis representing the running time. In one plot, the x-axis should represent the number of edges, and in the other, the number of nodes. If desired, include an average trend line across the data points.

Note:

If we are interested in testing different heuristic functions, we can repeat the above

experimental procedure with a different heuristic function. If we are interested in other performance metrics, such as memory usage or the number of nodes explored, we should also record them in Step 2 and plot them as the y-axis variable in additional graphs.

Effects of Arbitrary Heuristic Function

As an arbitrary heuristic is very unlikely to provide any meaningful estimate of the distance to the goal, it would not assist in guiding the search effectively in A_star. The algorithm heavily relies on the admissibility of the heuristic function to make informed decisions. Moreover, a poorly chosen or misleading heuristic can even degrade the performance of the algorithm. In such cases, A_star is no longer guaranteed to find the optimal path, whereas Dijkstra's algorithm still guarantees the shortest path since it does not rely on any heuristic.

Applications of A_star

A_star is particularly feasible in situations where we have an accurate heuristic function. For example, in navigation systems, it is impractical to use Dijkstra's algorithm due to the vast number of possible nodes in a map. However, A_star is much more suitable in this scenario, as it can derive its heuristic function from the straight-line (Euclidean) distance between a node and the destination (*dst*). In general, for applications where we are only interested in finding the shortest path between two specific nodes, and we have access to a reliable heuristic function, A_star is the preferred choice.

Part 5, Compare Shortest Path Algorithms

Our goal for the experiment is to compare two shortest path algorithms: A_star and Dijkstra. The experiment "graph" is the London Underground network, which is a real-life network with a complex structure, where lines serve as edges and stations serve as nodes. The data we will be using for the experiment is from the CSV files `london_stations.csv` and `london_connections.csv`.

We have the function `get_stations_data()` to obtain the stations' data from the CSV file, and `add_edges_from_csv()` will add each connection as an edge into the graph. Lastly, we have the `heuristic()` function to calculate the heuristic dictionary that will be used in A_star. For the heuristic values, we will use the Euclidean distance, treating latitude and longitude directly as x and y coordinates. This is not entirely accurate since the Earth is spherical, but as our graph is limited to the city of London, we can treat the space as flat for the purpose of this experiment.

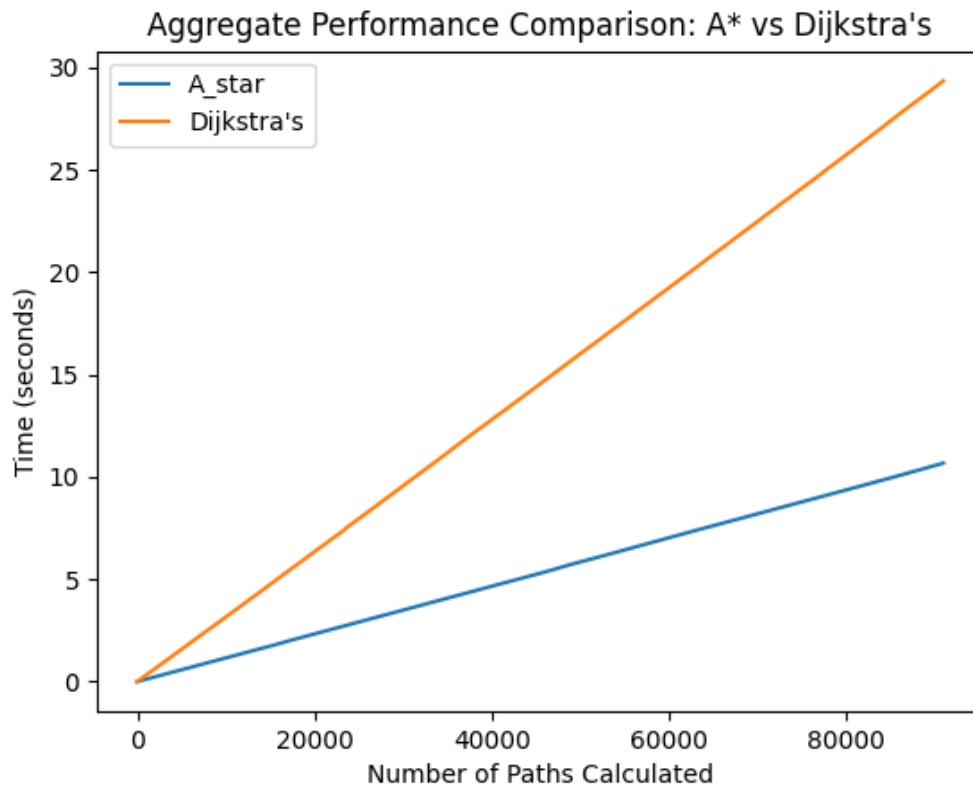


Figure 10: Aggregate Performance Comparison

After performing the experiment on both the A_star algorithm and Dijkstra's algorithm, we found that the A_star algorithm outperforms Dijkstra in runtime for combination of all stations. The Euclidean distance based on latitude and longitude is a decent choice of heuristic.




Figure 11: Performance Pre Iteration Comparison

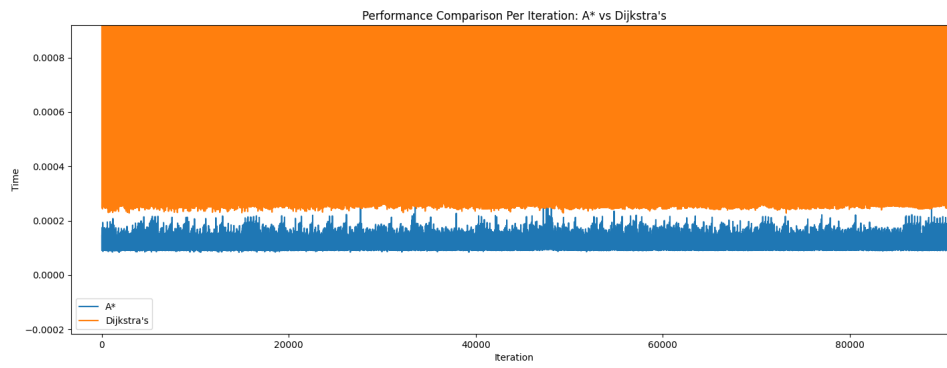


Figure 12: Performance Pre Iteration Comparison

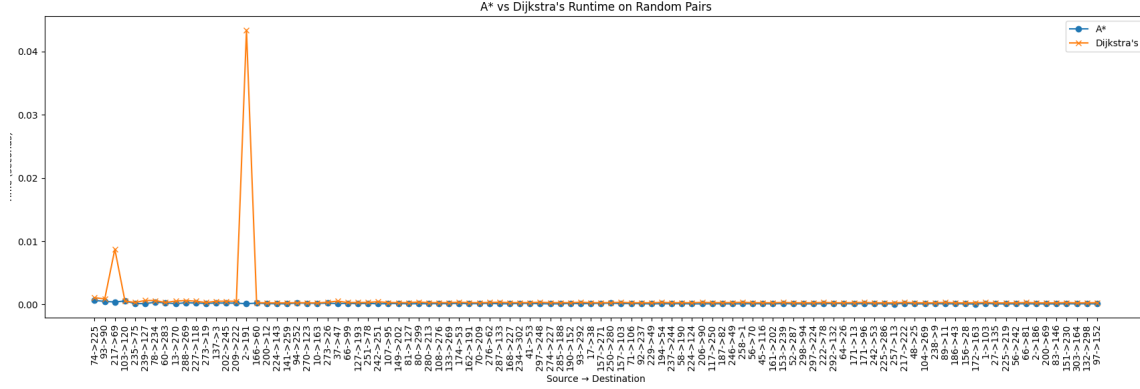


Figure 13: Performance Pre Iteration Comparison




Figure 14: Performance On random Path after Zoom in

However, there are still specific cases where Dijkstra slightly outperforms A_star. For example 103 \rightarrow 120 in the above, this may occur when the Euclidean distance calculated from the latitude and longitude between two stations is relatively small, but there is no direct line between them. To travel between these stations, we may need to take a detour, which makes the heuristic value much smaller than the actual cost, causing the A_star algorithm to struggle with performance.

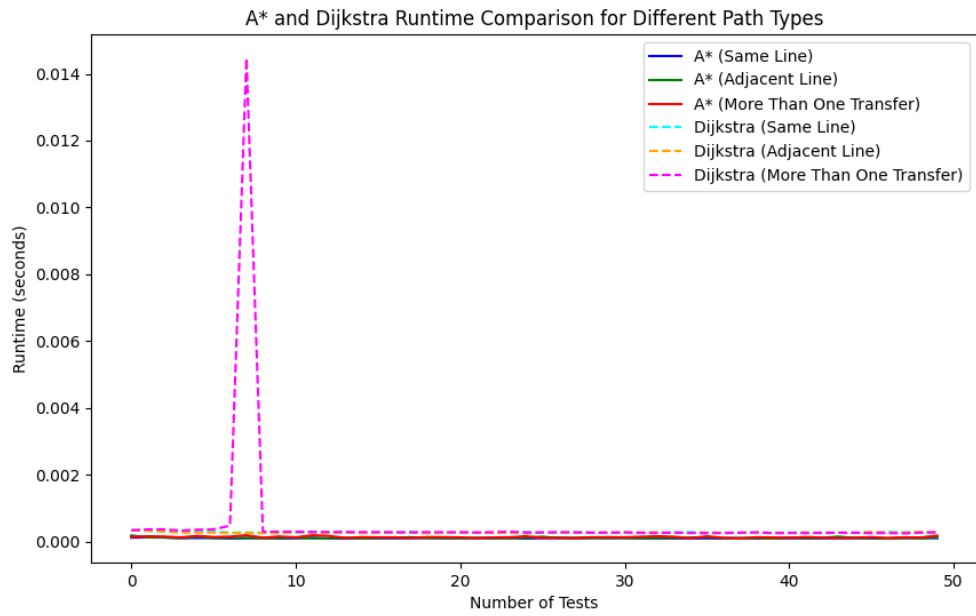


Figure 15: Performance On different number of lines

After zooming in:

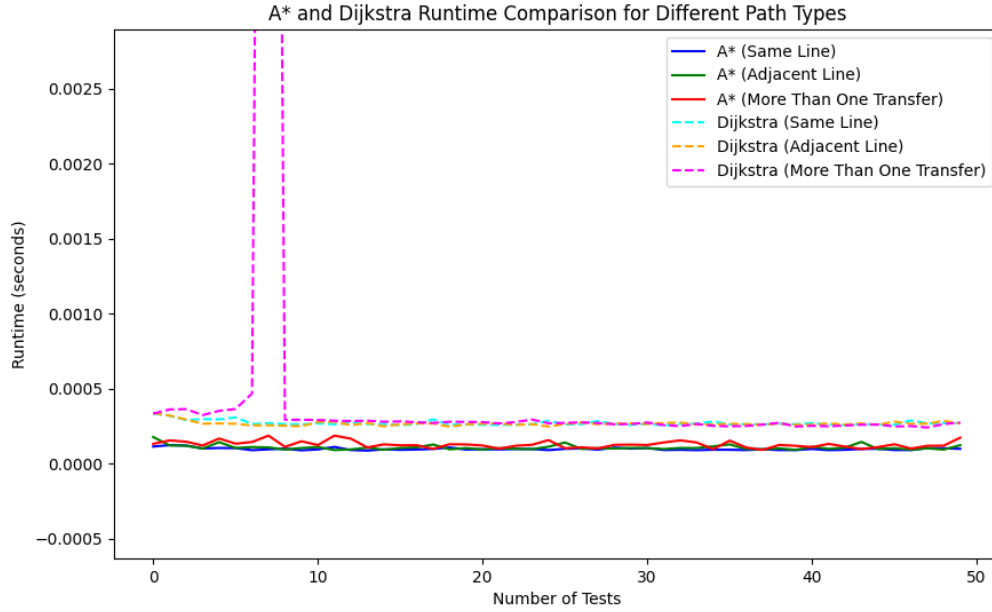


Figure 16: Performance On different number of lines after Zoom in

Upon further analysis, we observed that in scenarios involving stations on the same line, adjacent lines, or multiple lines, A_star still outperforms Dijkstra's algorithm in most cases. The rare cases where Dijkstra outperforms A_star can be attributed to the situation mentioned earlier: the heuristic value being much smaller than the actual cost due to the lack of a direct line between stations. This aligns with the idea that multiple transfer lines are necessary in such cases. In the above graph, we can see that in terms of the required transfers, more than one transfer lines are the closest together, except for the outliers when compared to any other pair of lines

When comparing these two algorithms, it is evident that A_star is more efficient in its routing capabilities than Dijkstra's in the London Underground network. The use of Euclidean distance in the heuristic function allows for more optimal searching. A_star performs better in terms of runtime, especially in cases where stations are located on the same line. Overall, from our experiment, we can conclude that A_star would be the best choice for routing in complex real-time network problems.

Part 6, Organize code per UML Diagram

The design principles and patterns being used in this diagram include composition, class inheritance, and adapters.

HeuristicGraph inherits from WeightedGraph, which inherits from Graph class.

Dijkstras, Bellman_Ford, and A_star (which is an adapter class for original A_star algorithm, another design principle) all inherit from SPAlgorithm.

Finally, ShortPathFinder uses composition (has-a relation) by passing a Graph object and interchangeable SPAlgorithm types into the class to help run. This way we can allow runtime switching of algorithms to whichever may be more efficient for the situation, especially using the set_algorithm function.

Currently, nodes are represented as only integers which is very limiting in that it cannot represent much information. Thus, a solution could be implementing a custom Node class, where we can define metadata for each node, like name, value, date, or something similar. A sample implementation could be:

```
class Node:
    def __init__(self, name: str, value: int, date: int):
        self.name = name
        self.value = value
        self.date = date

    def __repr__(self):
        return f"Node(name={self.name!r},
-----value={self.value},date={self.date})"
```

Now whenever we need a node, we can pass in this Node object, which contains a lot more data than just a simple integer value.

This means we must change all classes that utilize nodes. We can change adjacency list from List[int] to Dict[Node, List[Node]], or change weights to use Dict[Tuple[Node, Node], float]. Other sample changes include changing add_node(node: int) function, to add_node(node: Node). And when we need specific values from Node, we can simply call Node.name or Node.date, etc. This

would be a much better design for this UML.